

Xolaris LOE — 1-Month Go-Live Plan

Autonomous Security (Xolaris) · 11 September 2026 · 1 engineer · 22 working days (4 weeks)
Engines from the last 3 months stay as-is. This plan only covers hosting, database, and a thin live demo.
Hours are engineering time. GPU rent and Hugging Face download wait are extra.

22d This month (working days)	132h Engineering hours (6h/day)	4 wk Calendar	Out Wave D / nice-to-have
---	---	-------------------------	-------------------------------------

Why this is 22 days, not 83. The 83-day number was the full leftover backlog (RAG, real execute, extra scanners, OPA, full UI, Alembic, JWT). That work is *not* in this month. This plan ships a hosted demo: RavenX chat + live Nuclei/Trivy scan + findings stored in Postgres + a locked API. Engines are already written. We host them, we do not rebuild them.

1. What “done in 1 month” means

At the end of week 4 a reviewer can:

- 1. Open chat and get an answer from RavenX on the cloud VM (llama.cpp).
- 2. Run a live Nuclei or Trivy scan (not simulate-only).
- 3. See those findings and assets on GET /api/v1/findings and GET /api/v1/assets after a restart (Postgres, not SQLite).
- 4. Call the API with a key (not an open internet API).
- 5. After a plan is created, simulation runs and an approval request is submitted (no real patch on the server).

2. Four-week calendar

Week	Days	Hours	Work	Exit check
1 — Host	5	30	Rent Vast.ai (or RunPod) Linux GPU VM. Download RavenX from Hugging Face. Build llama.cpp on port 8080. Clone backend, pip install, .env, uvicorn.	/chat returns a RavenX answer
2 — Scanners + DB	5	30	Install Nuclei and Trivy (Prowler/Checkov/Grype only if time). Live scan smoke. Start Postgres 16 from docker-compose. Point DATABASE_URL. Confirm scan writes findings.	Live scan + GET /findings after restart
3 — Persist + test	6	36	Tenant isolation check. Save chat history in Postgres (replace in-memory MemoryService). E2E happy path: scan → findings → trust → risk → decision → plan.	Restart does not wipe findings or chat
4 — Demo lock	6	36	Simple API keys. Auto-run simulation and submit approval after a plan. Thin findings list in the existing React app (one page, real API). Rehearse a 10-minute demo.	Keyed API + demo script works
Total		22	132 hours	Hosted go-live demo

3. Day-by-day LOE (22 days)

ID	Item	Days	Hours	Cut from the 83-day plan
----	------	------	-------	--------------------------

W1.1	Cheap GPU Linux VM (Vast.ai first)	1.5	9	Was 2d
W1.2	Download RavenX from Hugging Face	0.5	3	Was 1d; download runs in parallel with VM setup
W1.3	Build llama.cpp and serve on :8080	1.5	9	Was 2d
W1.4	Backend clone, venv, requirements, .env, uvicorn	1.5	9	Was 1d; includes env wiring to llama.cpp
W2.1	Install Nuclei + Trivy on Linux (live PATH)	2	12	Was 3d; Prowler/Checkov/Grype deferred if time is short
W2.2	Smoke: live scan + chat + GET /findings	1	6	Same intent, tighter
W2.3	PostgreSQL (existing docker-compose) + DATABASE_URL	2	12	Was 3d; no Alembic this month (CREATE_TABLES)
W3.1	Multi-tenant check (tenant_id isolation + backup dump)	1.5	9	Was 3d; indexes/pool hardening later
W3.2	Persistent chat memory in Postgres	2	12	Was 3d
W3.3	E2E happy path only (scan through plan)	2.5	15	Was 5d; no full execute/verify/report matrix
W4.1	API keys (not full JWT/OIDC)	2	12	Was 5d
W4.2	Auto simulation + submit approval after plan	2.5	15	Was 4d; no auto-execute
W4.3	Findings list page on real API + demo rehearsal	1.5	9	Was 6d full UI; one screen only
Total this month		22	132	

4. Explicitly out of this month

Do not schedule these in the 22 days. They were most of the old 83-day total.

- Alembic production migrations
- JWT / OIDC / SSO
- Full React wiring (dashboard, pipeline, approvals UX)
- Async scan jobs (GET /scans/{id})
- Live threat feeds (CISA KEV, NVD) and threat-intel REST
- OPA policy transport
- OpenAI / Anthropic providers
- RAG / vector search
- Real execution (SSH/cloud patch) — keep the recording adapter
- Extra scanners: Gitleaks, ZAP, OSV, Semgrep; Prowler/Checkov/Grype live if week 2 slips
- MCP agent tools, nginx TLS polish, structured logging platform

5. Already built — do not rebuild

Scan/ingest, adapters (code), normalization, local policy, evidence, assets, findings/assets GET APIs (Feature 4.6), trust, risk, decision, planner, AI harness, simulation, approval, execution orchestration, verification, reporting (Feature 4.5.4), chat V2 to llama.cpp.

HTTP already mounted: POST /scans, GET /findings, GET /assets, trust, risk, decisions, plans, simulations, approvals, executions, verifications, reporting, ai-harness, /chat.

6. Microsoft Planner

One feature for this month: 4.7 One-Month Go-Live (22 days).

Keep 4.5 Reporting and 4.6 Findings/Assets complete.

Put the “out of this month” list in a later bucket (4.8 Backlog) with no dates until 4.7 is done.

Assumptions: 1 engineer, 6h/day, 5-day weeks, Vast.ai/RunPod GPU Linux, Nuclei+Trivy as the live scanners, Postgres 16, CREATE_TABLES not Alembic, API keys not SSO. If the GPU VM is delayed, week 1 slips and the month becomes chat-only on simulate scans — call that out immediately.